



Image filtering

Mårten Björkman

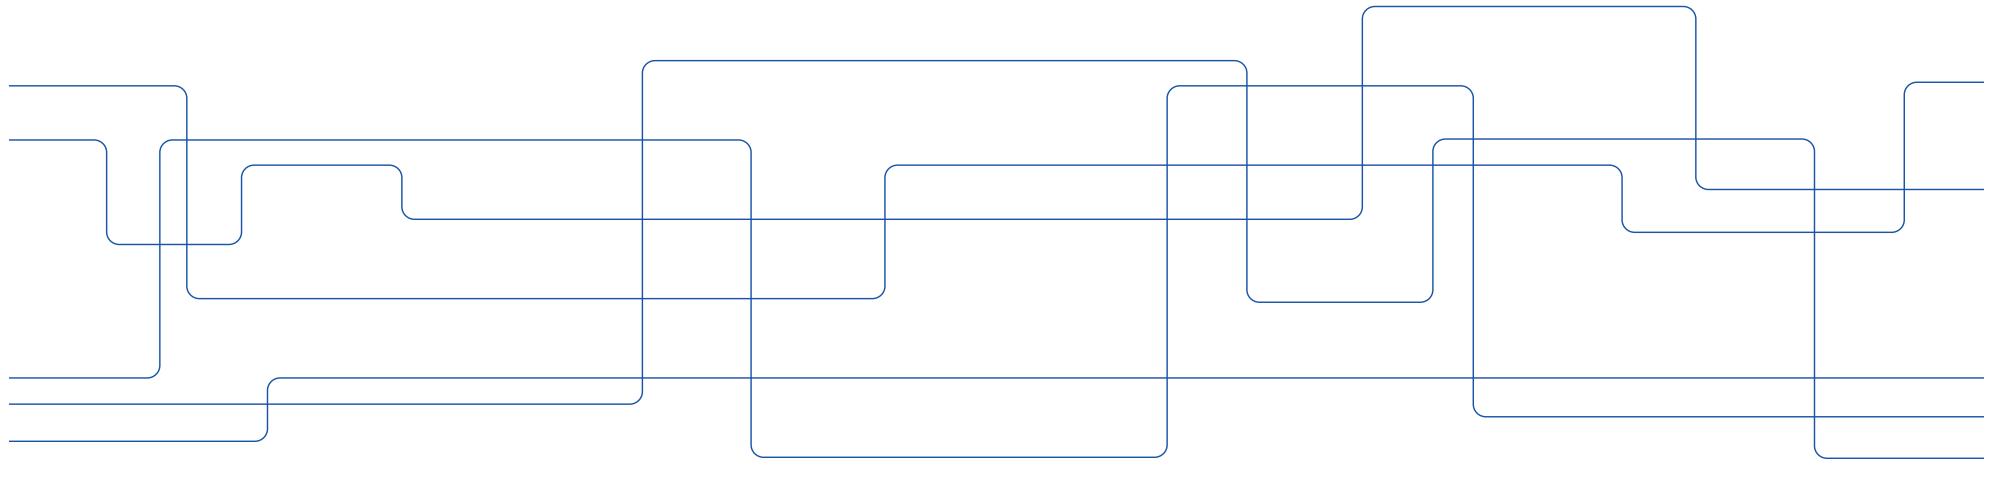
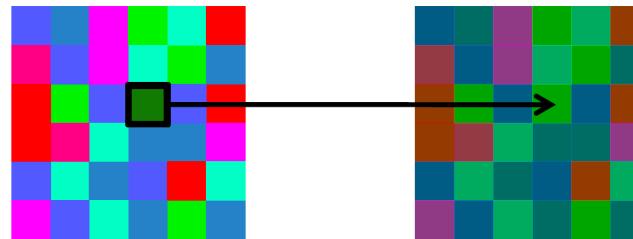




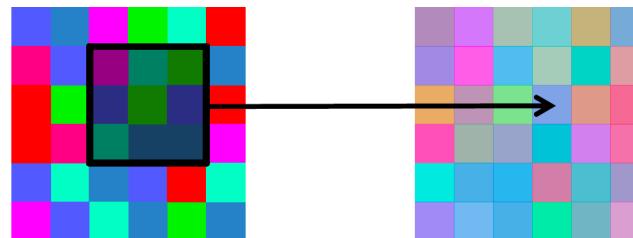
Image filtering

Point Operation



point processing

Neighborhood Operation



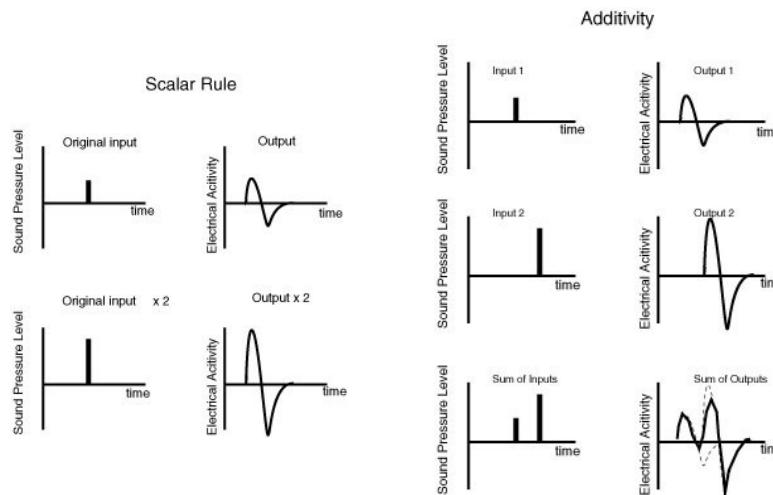
filtering

- Most spatial image filters work in local neighbourhoods. You combine a group of pixels in a neighbourhood to compute a new value for the pixel in the centre.
- Most filters are also Linear Shift-Invariant, that are easy to implement and analyse.



Linear operators

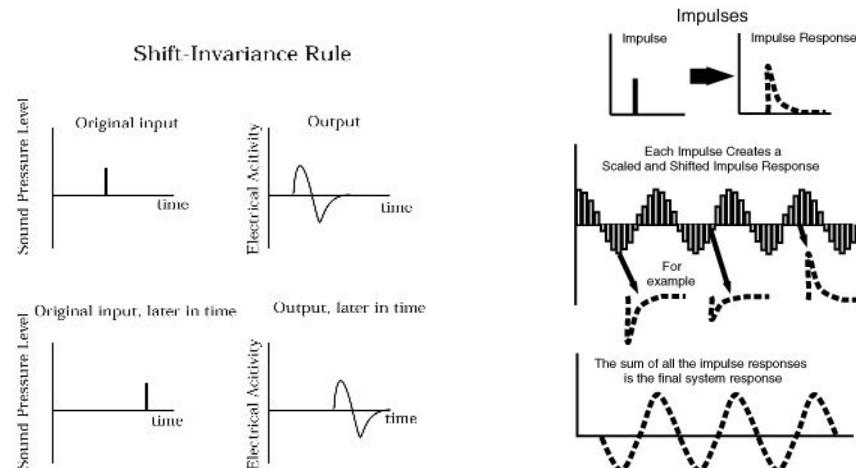
- Linear systems theory can be used to understand linear operators. A linear operator obeys the principle of superposition:
 - Homogeneity (scalar rule): an increase in strength of the input, increases the output/response for the same amount.
 - Additivity: if the input consists of two signals, the output/response is equal to the sum of the individual responses.





Linear shift-invariant operators

- Shift-invariance: If a system is given two impulses with a time delay, the response remains the same except for a time difference.
 - Signals can be represented as sums of impulses of different strengths (image intensities), shifted in time (image space).
- If we know how system responds to an impulse, we know how it reacts to all combinations of impulses: impulse-response function.





Linear filters in image processing

Digital linear filters modify pixel values based on some neighbourhoods. Linear means the output is a weighted combination of neighbouring pixels.

$$g(x, y) = \sum_{i,j} w(i, j)f(x + i, y + j)$$

Shift-invariance here means that the weights $w(i, j)$ are the same for all pixels.

- Benefits:
 - Linear filters are simpler and easier to understand than non-linear filters.
 - Can combine multiple linear filters applied in sequence into a single linear filter.
 - Sometimes linear filters are easier to invert, at least theoretically.



Applying a linear filter

input

f_{11}	f_{12}	f_{13}	f_{14}	f_{15}
f_{21}	f_{22}	f_{23}	f_{24}	f_{25}
f_{31}	f_{32}	f_{33}	f_{34}	f_{35}
f_{41}	f_{42}	f_{43}	f_{44}	f_{45}
f_{51}	f_{52}	f_{53}	f_{54}	f_{55}

filter

h_{11}	h_{12}	h_{13}
h_{21}	h_{22}	h_{23}
h_{31}	h_{32}	h_{33}

output

We have a 5×5 pixel image and a 3×3 point filter that we sweep over the image.



Applying a linear filter

input

hf_{11}	hf_{12}	hf_{13}	f_{14}	f_{15}
hf_{21}	hf_{22}	hf_{23}	f_{24}	f_{25}
hf_{31}	hf_{32}	hf_{33}	f_{34}	f_{35}
f_{41}	f_{42}	f_{43}	f_{44}	f_{45}
f_{51}	f_{52}	f_{53}	f_{54}	f_{55}

filter

h_{11}	h_{12}	h_{13}
h_{21}	h_{22}	h_{23}
h_{31}	h_{32}	h_{33}

output

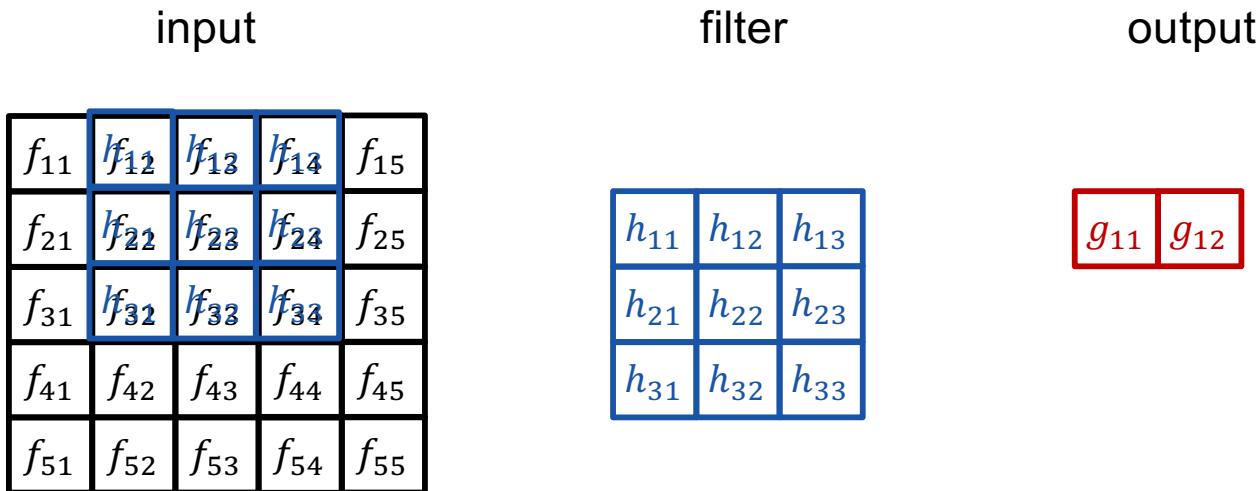
g_{11}

$$g_{11} = h_{11}f_{11} + h_{12}f_{12} + h_{13}f_{13} + h_{21}f_{21} + h_{22}f_{22} + \dots + h_{33}f_{33}$$

At the first position we do a point-wise multiplication and then sum up.



Applying a linear filter



$$g_{12} = h_{11}f_{12} + h_{12}f_{13} + h_{13}f_{14} + h_{21}f_{22} + h_{22}f_{23} + \dots + h_{33}f_{34}$$

At the second position we do a point-wise multiplication and then sum up.



Applying a linear filter

input

f_{11}	f_{12}	f_{13}	f_{14}	f_{15}
f_{21}	f_{22}	f_{23}	f_{24}	f_{25}
f_{31}	f_{32}	f_{33}	f_{34}	f_{35}
f_{41}	f_{42}	f_{43}	f_{44}	f_{45}
f_{51}	f_{52}	f_{53}	f_{54}	f_{55}

filter

h_{11}	h_{12}	h_{13}
h_{21}	h_{22}	h_{23}
h_{31}	h_{32}	h_{33}

output

g_{11}	g_{12}	g_{13}
----------	----------	----------

$$g_{13} = h_{11}f_{13} + h_{12}f_{14} + h_{13}f_{15} + h_{21}f_{23} + h_{22}f_{24} + \dots + h_{33}f_{35}$$

At the third position we do a point-wise multiplication and then sum up.



Cross-correlation vs Convolution

So far we have used a cross-correlation

$$g(x, y) = \sum_{i,j} h(i, j)f(x + i, y + j) \quad g = h \otimes f$$

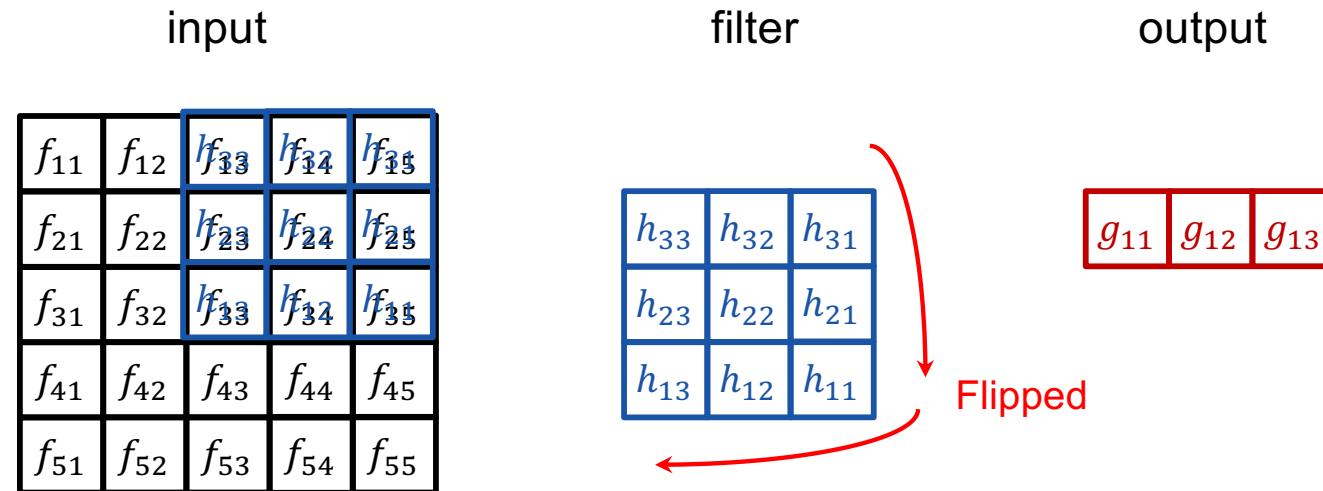
But usually we do convolutions instead

$$g(x, y) = \sum_{i,j} h(i, j)f(x - i, y - j) \quad g = h * f$$

Reasons:

- Linear system theory tells us that any linear shift-invariant (LSI) operator can be written as a convolution.
- Convolutions are commutative, e.g. $g = h * f = f * h$, so it doesn't matter in which order you apply the filters.

Applying a linear filter with a convolution



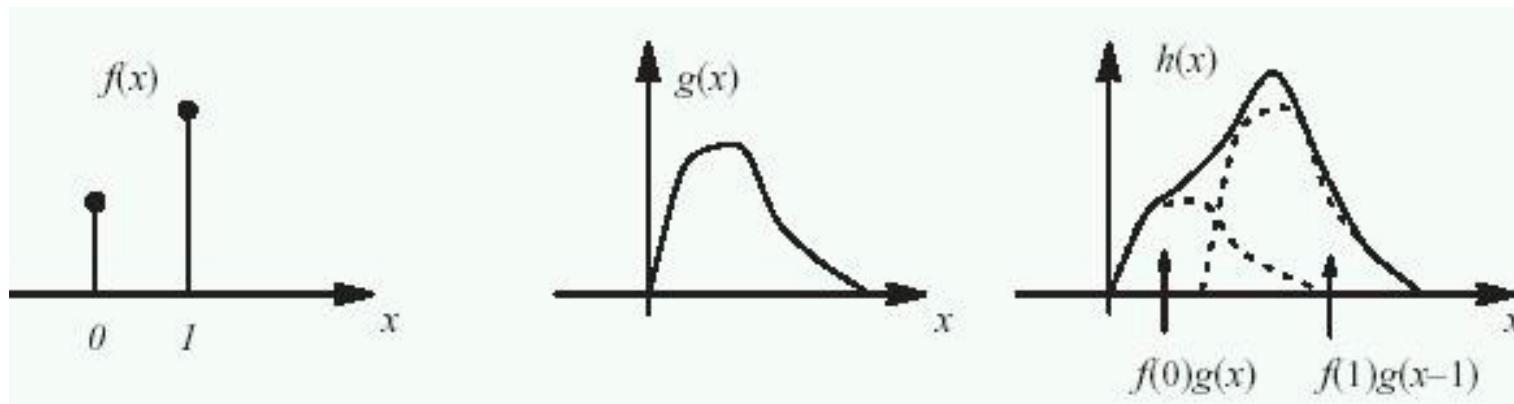
$$g_{13} = h_{33}f_{13} + h_{32}f_{14} + h_{31}f_{15} + h_{23}f_{23} + h_{22}f_{24} + \dots + h_{11}f_{35}$$

Essentially the same, but with filter kernel flipped x-wise and y-wise.

Understanding convolutions

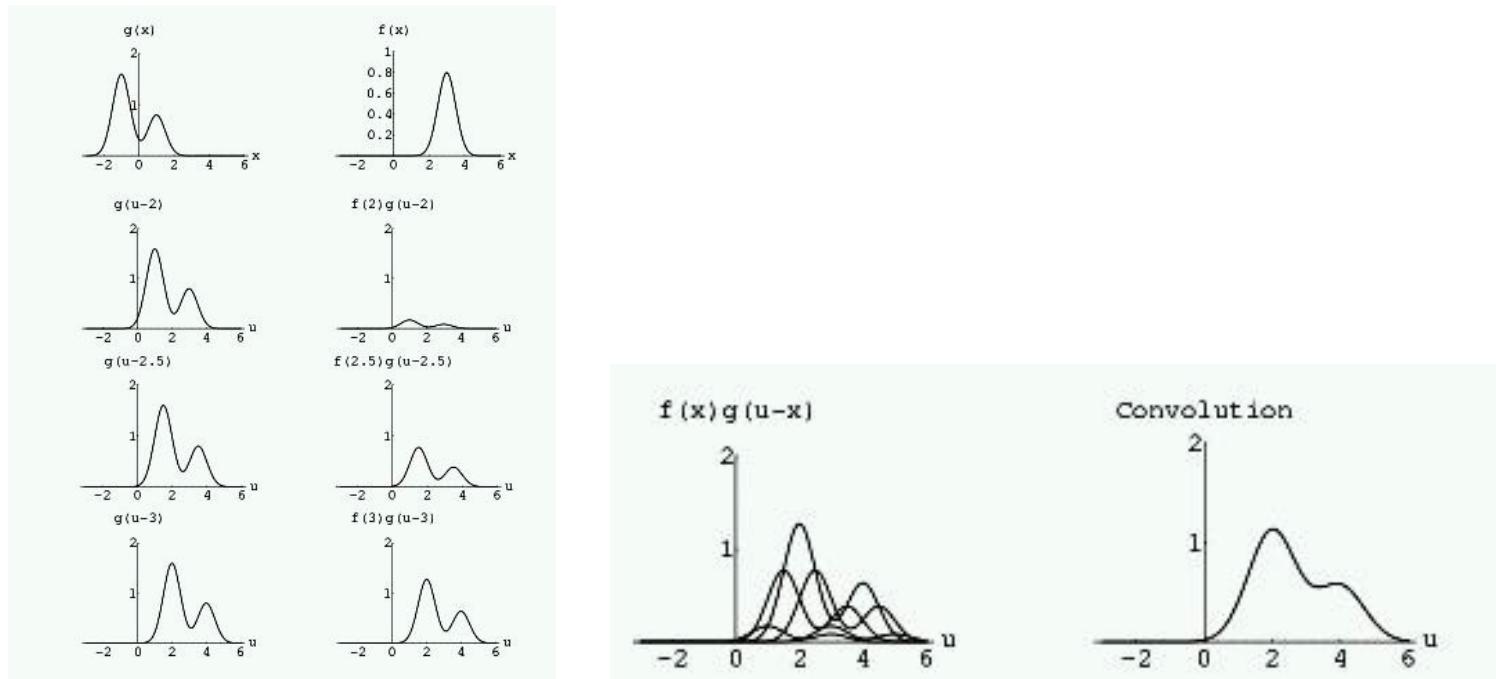
- Mathematically, a convolution is defined as the integral over space of one function at α , times another function at $x - \alpha$.

$$f(x) * g(x) = \int_{-\infty}^{\infty} f(\alpha)g(x - \alpha)d\alpha = g(x) * f(x) = \int_{-\infty}^{\infty} g(\alpha)f(x - \alpha)d\alpha$$



Understanding convolutions

- Way of considering convolution: weighted sum of shifted copies of one function, with weights given by the function value of the second function.





The filter kernel

$$g(x, y) = \sum_{i,j} h(i, j) f(x - i, y - j)$$

h_{11}	h_{12}	h_{13}
h_{21}	h_{22}	h_{23}
h_{31}	h_{32}	h_{33}

- In the discrete case, the filter is represented by an array of weights.
- It goes by many names:
 - Impulse response
 - Point spread function
 - Filter kernel
 - Filter mask
 - Template



Convolution: 1D example

If

$$F_1 = [1, 2, 3, 4, 5]$$

$$F_2 = [1, 2, 1, 2, 1]$$

$$G_1 = [-1, 2, -1]$$

$$G_2 = [1, 2, 3]$$

then

$$F_1 * G_1 = [-1, 0, 0, 0, 0, 6, -5]$$

$$F_2 * G_1 = [-1, 0, 2, -2, 2, 0, -1]$$

$$F_1 * G_2 = [1, 4, 10, 16, 22, 22, 15]$$

$$F_2 * G_2 = [1, 4, 8, 10, 8, 8, 3]$$

- Note1: outside the windows, values are usually assumed to be zero.
- Note2: normally you assume $x = 0$ at centre of filter kernel.



Convolution: 1D example

$$F_1 = [1, 2, 3, 4, 5]$$

$$G_2 = [1, 2, 3]$$

$$\begin{array}{r} & 1 & 2 & 3 & 4 & 5 \\ * & & & 1 & 2 & 3 \\ \hline & 3 & 6 & 9 & 12 & 15 \\ & 2 & 4 & 6 & 8 & 10 \\ + & 1 & 2 & 3 & 4 & 5 \\ \hline & 1 & 4 & 10 & 16 & 22 & 22 & 15 \end{array}$$

An easier way of doing it! Almost like regular multiplication.



Convolutions in Matlab and Python

In Matlab (and SciPy.signal) there are essentially two functions that can be used for image filtering.

- `conv2` (`convolve2d`) – a proper convolution according to the theory.

$$g(x) = \sum_i h(i) f(x - i)$$

Use it when you want to stick as close as possible to the theory and exploit known mathematical properties of convolutions.

- `filter2` (`correlate2d`) – does not flip the kernel before applying it.

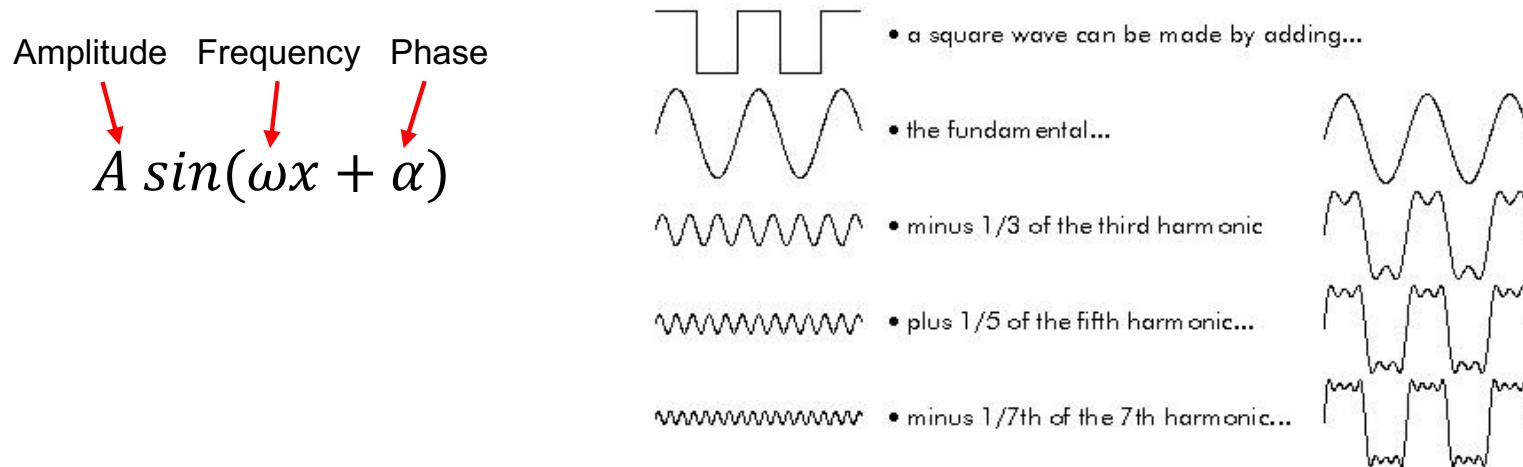
$$g(x) = \sum_i h(i) f(x + i)$$

Use it when you want to match a known shape to the image data to see where in the image you have the best correlation.



Signal decomposition

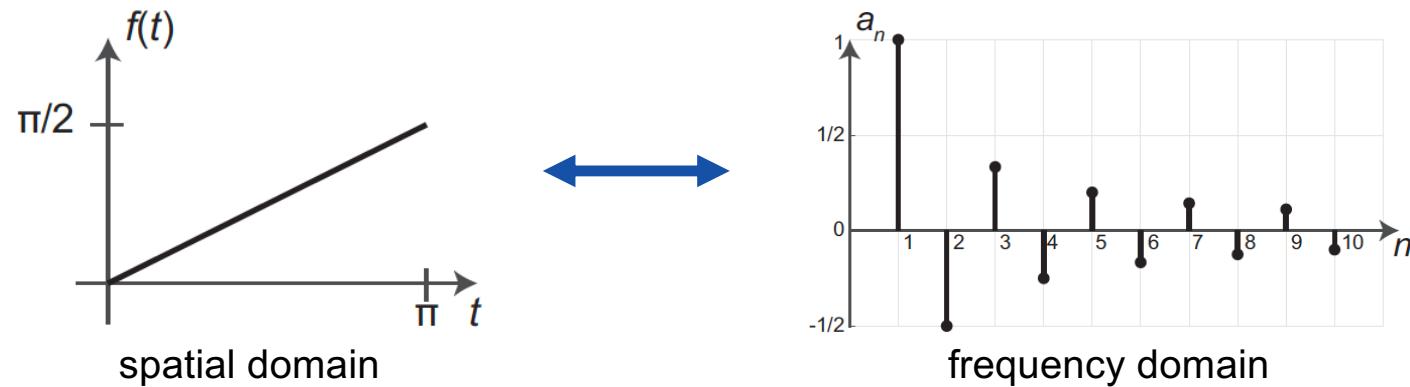
- We can think of an image as a weighted sum of pixels, but we can think of it as a weighted sum of something else.
- In 1807 Jean Baptiste Fourier showed that any periodic signal can be represented by a sum of sine waves with appropriate amplitude, frequency and phase.





Fourier series

$$\frac{1}{2}t = \sin(t) - \frac{1}{2}\sin(2t) + \frac{1}{3}\sin(3t) - \frac{1}{4}\sin(4t) + \dots$$

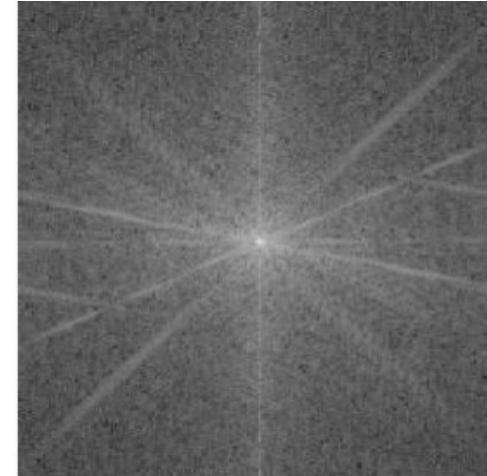


- By writing the signal (image) as a sum of sinusoids, you get an alternative frequency space representation of the same signal.



Fourier Transform

- The Fourier transform is an operation that calculates the frequency, amplitude and phase of each sine wave needed to make up any given signal.
- In the Fourier space, each point represents a particular frequency contained in the original spatial domain image.
- The Fourier Transform is used in a wide range of applications, such as image analysis, image filtering, and image compression.





Fourier Transform in 2D

$$\mathcal{F}(f(x)) = \int_{x \in \mathbb{R}^2} f(x) e^{-i\omega^T x} dx = \hat{f}(\omega)$$

$$\mathcal{F}^{-1}(\hat{f}(\omega)) = \frac{1}{(2\pi)^2} \int_{\omega \in \mathbb{R}^2} \hat{f}(\omega) e^{i\omega^T x} d\omega = f(x)$$

$$e^{i\omega^T x} = \cos \omega^T x + i \sin \omega^T x, \quad \omega^T x = \omega_1 x_1 + \omega_2 x_2$$

Terminology:

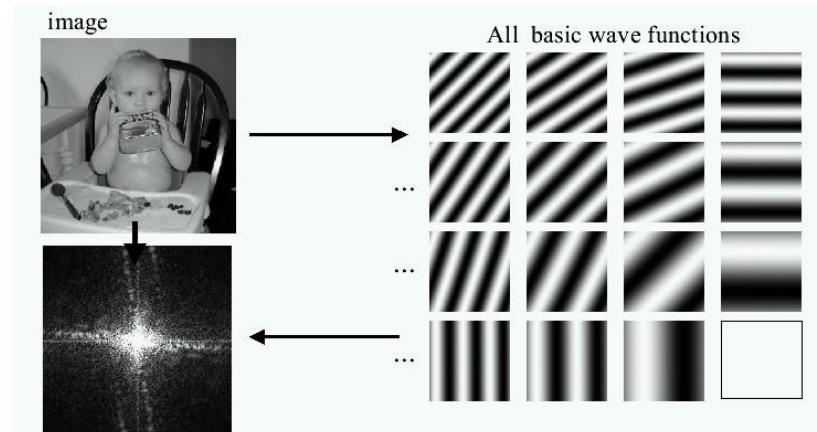
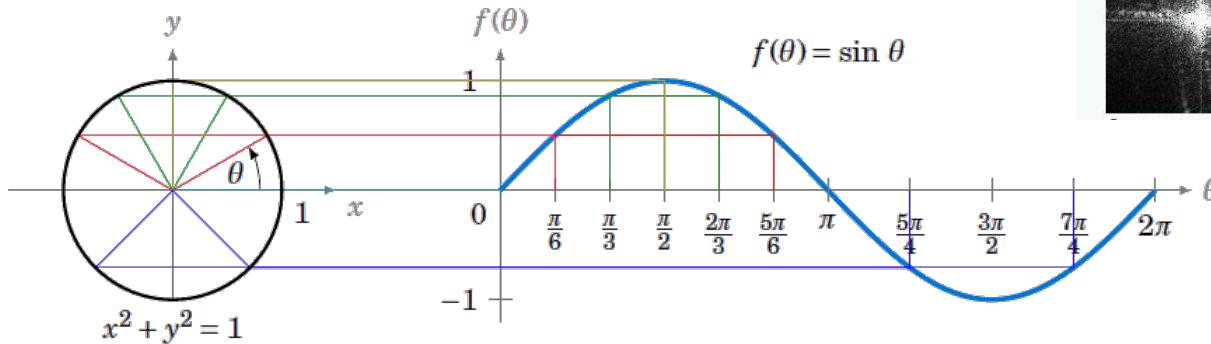
- Frequency spectrum: $\hat{f}(\omega) = Re[\hat{f}(\omega)] + i Im[\hat{f}(\omega)] = |\hat{f}(\omega)| e^{-i\phi(\omega)}$
- Angular frequency: $\omega = \begin{pmatrix} \omega_1 \\ \omega_2 \end{pmatrix}$
- Frequency: $f = \frac{\omega}{2\pi}$
- Wavelength: $\lambda = \frac{2\pi}{||\omega||} = \frac{2\pi}{\sqrt{\omega_1^2 + \omega_2^2}}$



Fourier Transform in 2D

$$Re[\hat{f}(\omega)] = \int_x f(x) \cos(\omega^T x) dx$$

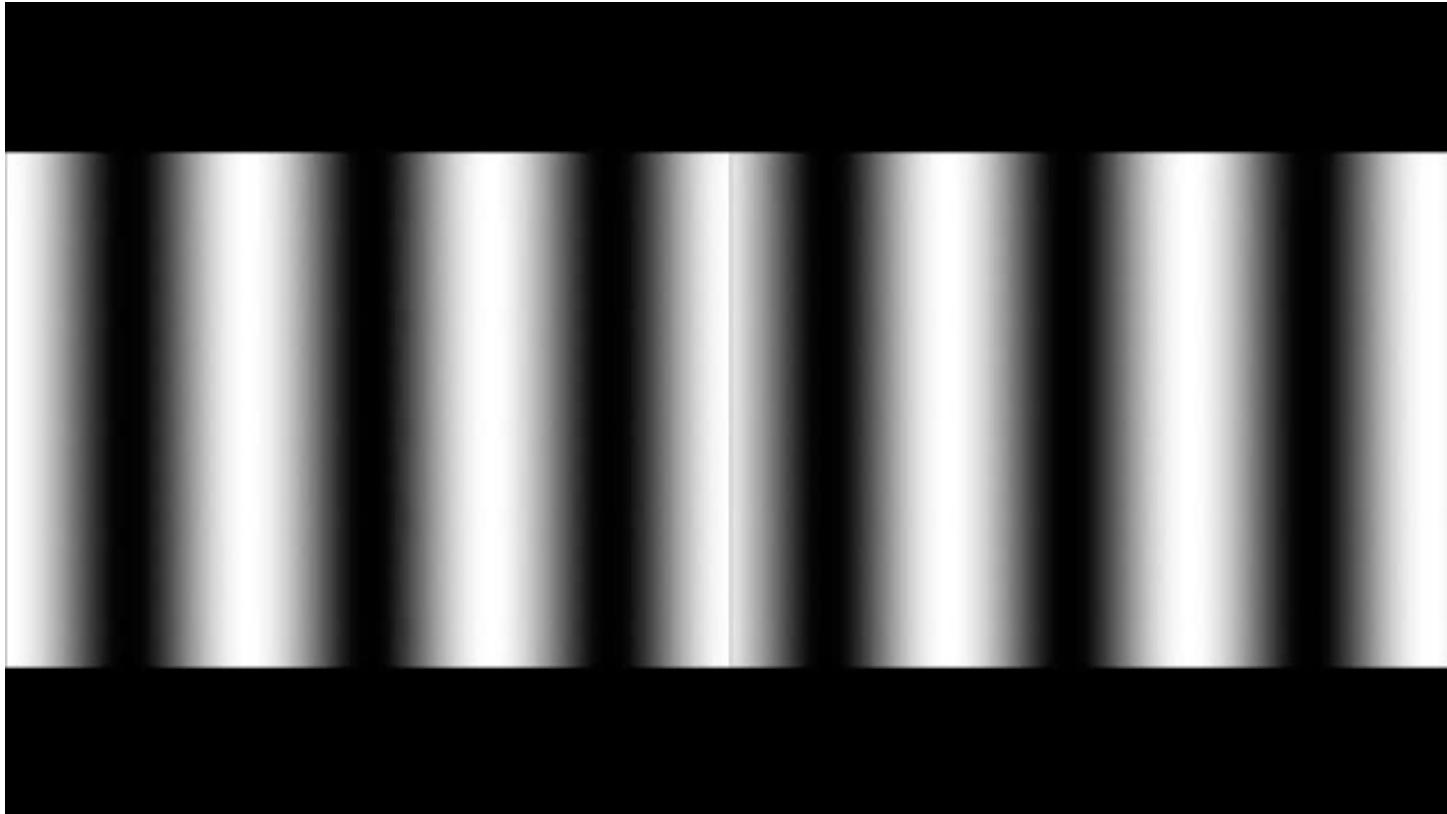
$$Im[\hat{f}(\omega)] = - \int_x f(x) \sin(\omega^T x) dx$$



- The complex exponential function $e^{-i\omega^T x}$ can be seen as rotating around the unit circle.
- You correlate $f(x)$ twice, first with $\cos(\omega^T w)$ and then with $\sin(\omega^T w)$.



Image decomposition



Gradually reconstruct image by summing up waveforms in order of decreasing magnitudes



Discrete Fourier Transform (DFT)

$$\hat{f}(u, v) = \sum_{x=0}^{M-1} \sum_{y=0}^{N-1} f(x, y) e^{-2\pi i (\frac{xu}{M} + \frac{yu}{N})}$$

$$f(x, y) = \frac{1}{MN} \sum_{u=0}^{M-1} \sum_{v=0}^{N-1} \hat{f}(u, v) e^{+2\pi i (\frac{xu}{M} + \frac{yu}{N})}$$

Terminology:

- Frequency spectrum: $\hat{f}(u, v) = Re[\hat{f}(u, v)] + i Im[\hat{f}(u, v)] = |\hat{f}(u, v)| e^{-i\phi(u, v)}$
- Power spectrum: $P(u, v) = |\hat{f}(u, v)|^2 = Re[\hat{f}(u, v)]^2 + Im[\hat{f}(u, v)]^2$

The magnitude $|\hat{f}(u, v)|$ is simply the peak value, and the phase $\phi(u, v)$ determines where the origin is, or where the sinusoid starts.

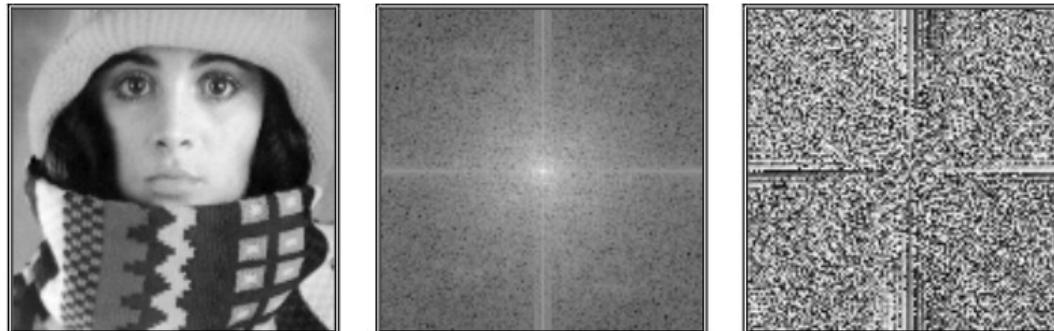


Magnitude and Phase

- The Fourier coefficients $\hat{f}(u, v)$ are complex numbers, but it is not obvious what the real and imaginary parts represent.
- Another way to represent the data is with phase and magnitude.

Magnitude: $|\hat{f}(u, v)| = \sqrt{Re^2[\hat{f}(u, v)] + Im^2[\hat{f}(u, v)]}$

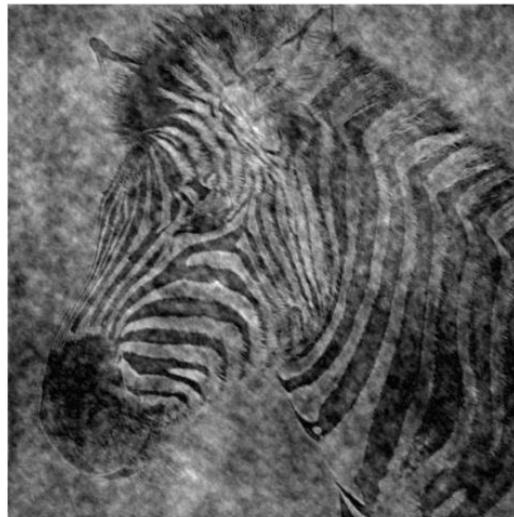
Phase: $\phi(u, v) = \tan^{-1} \frac{Im[\hat{f}(u, v)]}{Re[\hat{f}(u, v)]}$



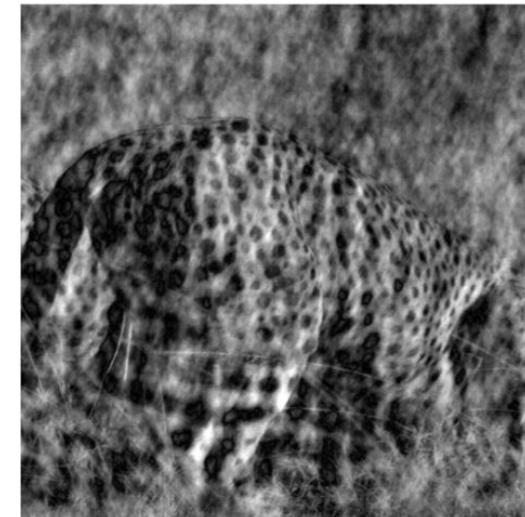


Magnitude and Phase

- Phase defines how waveforms are shifted along its direction.
 - Where edges will end up in the image (most important).
- Magnitude defines how large the waveforms are.
 - What grey-levels are on either side of edges (less important).



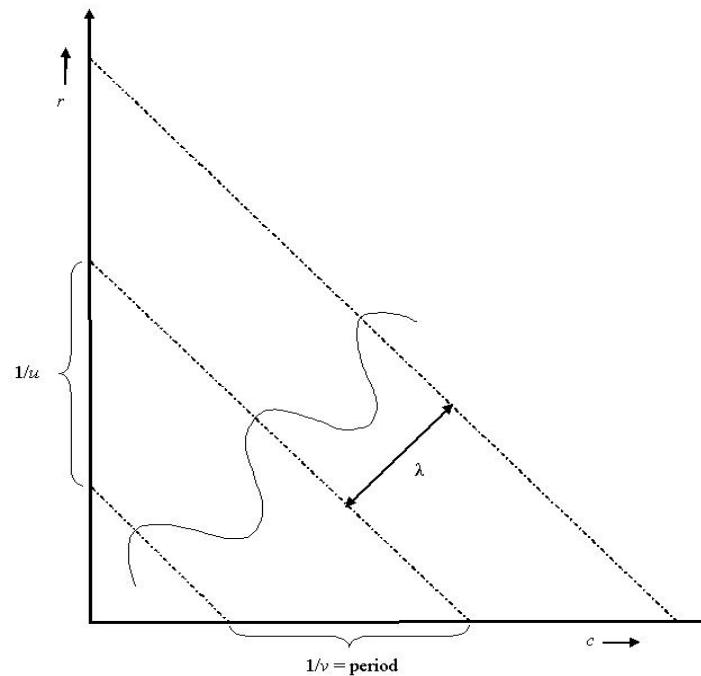
Phase of zebra - magnitude of tiger



Phase of tiger - magnitude of zebra

Wavelength

- The wavelength of the sinusoid is: $\lambda = \frac{1}{\sqrt{u^2+v^2}}$, where (u, v) are the frequencies along (x, y) and the periods are $1/u$ and $1/v$.





Change of basis functions

- An image can be viewed as a spatial array of grey-level values, but can also be thought of as a spatially varying function.
- It is possible to decompose the image into a set of orthogonal basis functions.
- When basis functions are combined (linearly) the original function will be reconstructed.
 - Spatial domain: basis consists of Dirac pulses, one per pixel.
 - Fourier domain: basis consists of complex exponential functions.
- The Fourier transform is “just” a change of basis functions!



Change of basis functions

- Assume you have a vector

$$\begin{pmatrix} 1 \\ 2 \\ 3 \\ 4 \end{pmatrix} = 1 \begin{pmatrix} 1 \\ 0 \\ 0 \\ 0 \end{pmatrix} + 2 \begin{pmatrix} 0 \\ 1 \\ 0 \\ 0 \end{pmatrix} + 3 \begin{pmatrix} 0 \\ 0 \\ 1 \\ 0 \end{pmatrix} + 4 \begin{pmatrix} 0 \\ 0 \\ 0 \\ 1 \end{pmatrix}$$

- This may be expressed with another basis (e.g. Haar wavelet).

$$\begin{pmatrix} 1 \\ 2 \\ 3 \\ 4 \end{pmatrix} = 2.5 \begin{pmatrix} 1 \\ 1 \\ 1 \\ 1 \end{pmatrix} - 1 \begin{pmatrix} 1 \\ 1 \\ -1 \\ -1 \end{pmatrix} - 0.5 \begin{pmatrix} 1 \\ -1 \\ 0 \\ 0 \end{pmatrix} - 0.5 \begin{pmatrix} 0 \\ 0 \\ 1 \\ -1 \end{pmatrix}$$

- The only condition is that the basis vectors are orthogonal.

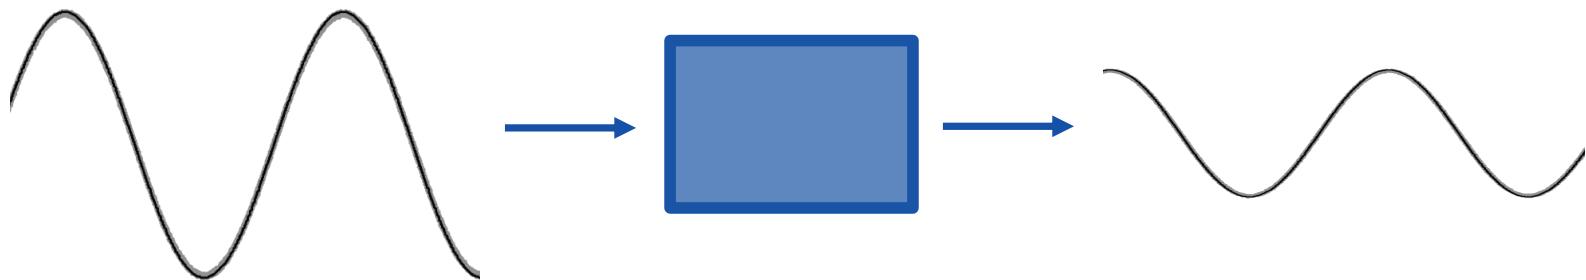


Why Fourier transforms?

- Complex exponential functions are eigenfunctions of convolutions!

$$e^{i\omega t} \implies A(\omega)e^{i\omega t}$$

- Note: $A(\omega)$ is complex (change in magnitude and phase).
- We can understand what the filter does to the different frequencies of the image.





The Convolution Theorem

- Convolution in the spatial domain is the same as multiplication in the Fourier (frequency) domain.

$$\mathcal{F}(h * f) = \mathcal{F}(h)\mathcal{F}(f)$$

$$f \rightarrow \boxed{*h} \rightarrow g = h * f \quad \hat{f} \rightarrow \boxed{\hat{h}} \rightarrow \hat{g} = \hat{h} \hat{f}$$

- (vice versa) Multiplication in the spatial domain is the same as convolution in the Fourier (frequency) domain.

$$\mathcal{F}(hf) = \mathcal{F}(h) * \mathcal{F}(f)$$

- Usage:
 - For analysis and understanding of linear shift-invariant filters.
 - Some filters are more easily represented in the Fourier domain.
 - Implementation: when the size of the filter is too large, it is more effective to use multiplication in the Fourier domain.



Transfer functions

- A linear, shift-invariant filter $h(x, y)$ is completely specified by its response to an impulse, which is called the impulse response.
- Convolution with an impulse function

$$h(x, y) * \delta(x, y) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} h(x_0, y_0) \delta(x - x_0, y - y_0) dx_0 dy_0 = h(x, y)$$

results in a "copy" of $h(x, y)$ to the location of the impulse.

- The transfer function $H(u, v)$ is the Fourier transform of the impulse response.
- Using the convolution theorem to describe the effects of the system:

$$g(x, y) = h(x, y) * f(x, y)$$

$$G(u, v) = H(u, v)F(u, v)$$



Spatial separability

- Assume a filter can be written as

$$\begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} = \begin{bmatrix} 1 \\ 0 \\ -1 \end{bmatrix} [1 \ 2 \ 1]$$

$$h(x, y) = h_1(x)h_2(y)$$

$$(h^*f)(x, y) = \sum_{i=0}^m \sum_{j=0}^m h(i, j)f(x - i, y - j) = \sum_{i=0}^m h_1(i) \left(\sum_{j=0}^m h_2(j)f(x - i, y - j) \right)$$

- If convolution mask $h(x, y)$ can be separated as above, 2D convolution can be performed as a series of 1D convolutions.
- Discrete case: If the mask is m^2 in size $\rightarrow 2m$ operations per pixel, instead of m^2 operations per pixel.



Separability in the Fourier domain

Fourier transforms are always separable

$$\hat{f}(u, v) = \sum_{x=0}^{M-1} \sum_{y=0}^{N-1} f(x, y) e^{-2\pi i (\frac{xu}{M} + \frac{yv}{N})} = \sum_{x=0}^{M-1} \left(\sum_{y=0}^{N-1} f(x, y) e^{-2\pi i \frac{yu}{M}} \right) e^{-2\pi i \frac{yv}{N}}$$

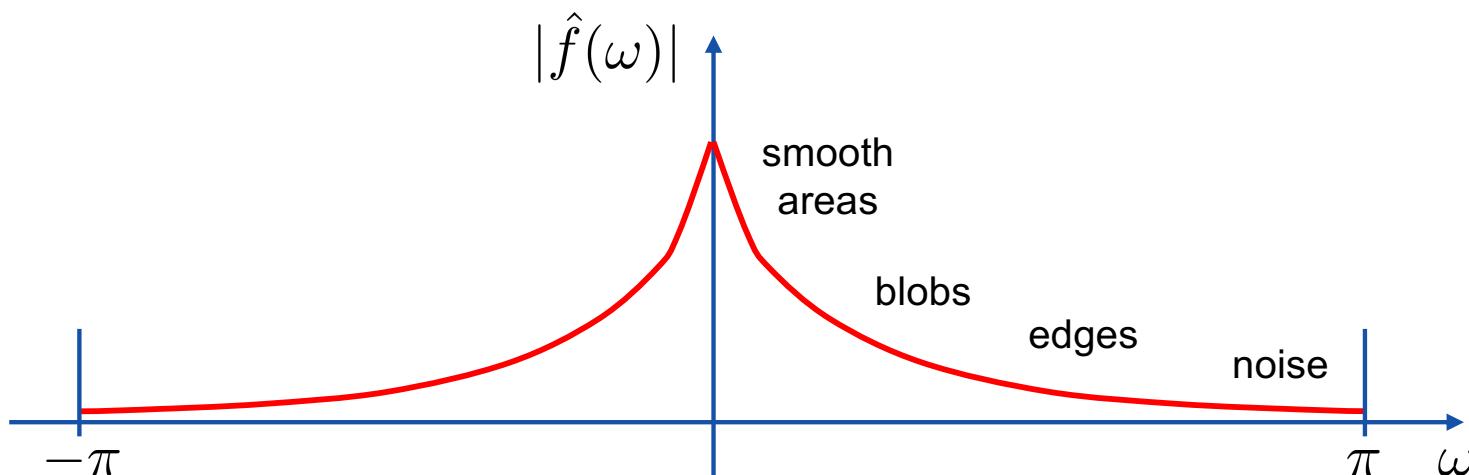
- A Fourier transform in 2D can always be performed as a series of two 1D Fourier transforms.
- Implication: For large non-separable spatial filters, it is faster to
 1. Go to the frequency domain (Fourier transform)
 2. Apply the filter as a point-wise multiplication
 3. Go back to the spatial domain (Inverse Fourier transform)



Applications

Filtering techniques typically modify frequency characteristics:

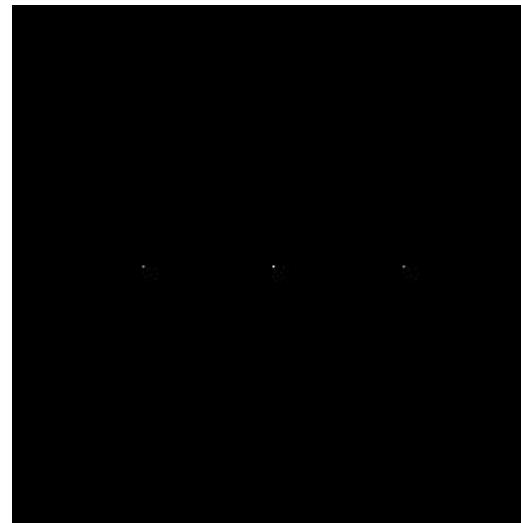
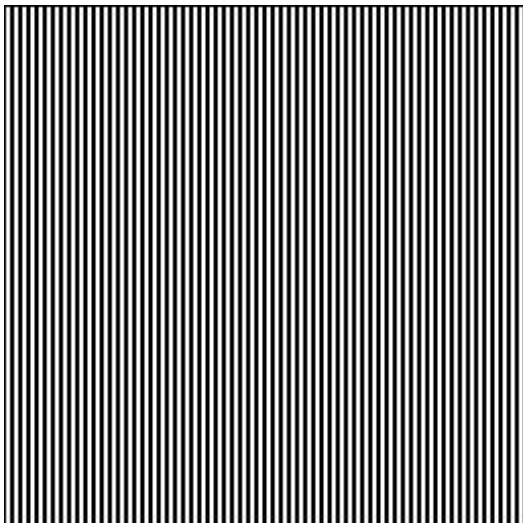
- Remove noise (decrease high frequencies)
- Smooth (decrease high frequencies, increase low frequencies)
- Enhance edges (increase medium frequencies)





The maximum frequency

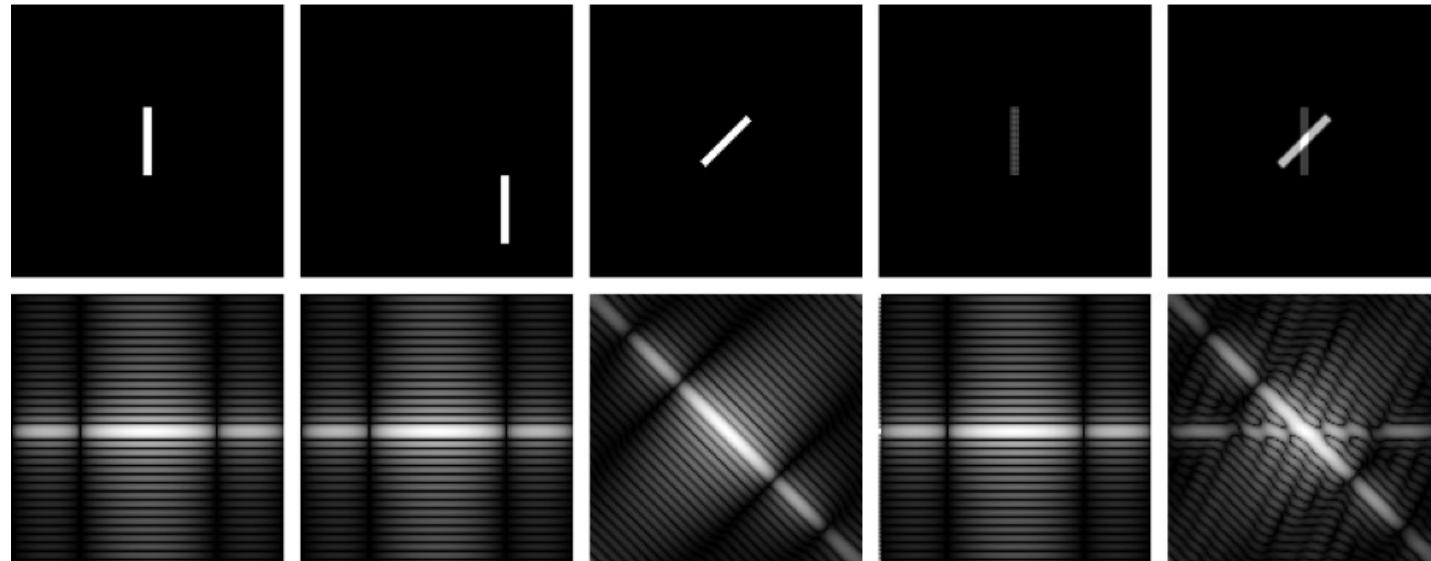
- The maximum frequency which can be represented in the spatial domain are one pixel wide stripes (period=2): $\omega_{max} = 2\pi \frac{1}{2} = \pi$
- So, 2 pixel wide stripes (period=4) give $\omega = 2\pi \frac{1}{4} = \frac{1}{2}\omega_{max}$



- Two points halfway between centre and the edge of the image, i.e. the represented frequency is half of the maximum.
- One point in the middle shows the DC-value (image mean).



Example: Transformations

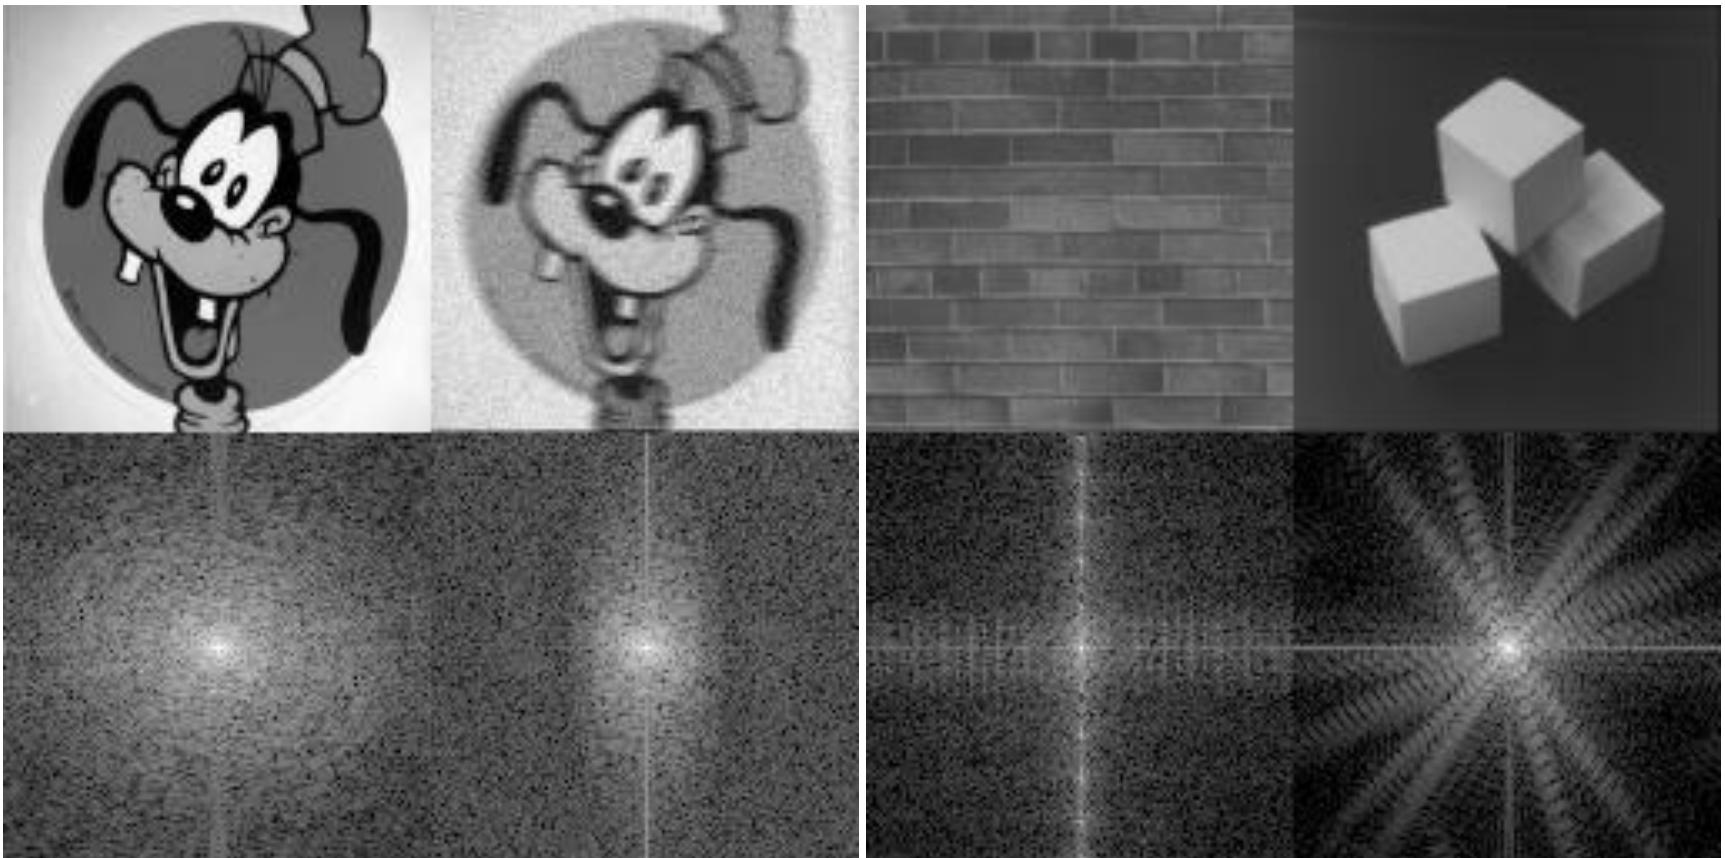


Observations:

- Translation only affects phase (not shown), not magnitude (shown)
- A rotation in one domain becomes a rotation in the other domain
- Fourier transform is a linear operation, $\mathcal{F}(af(x) + bg(x)) = a\mathcal{F}(f(x)) + b\mathcal{F}(g(x))$



Example images and Fourier transforms



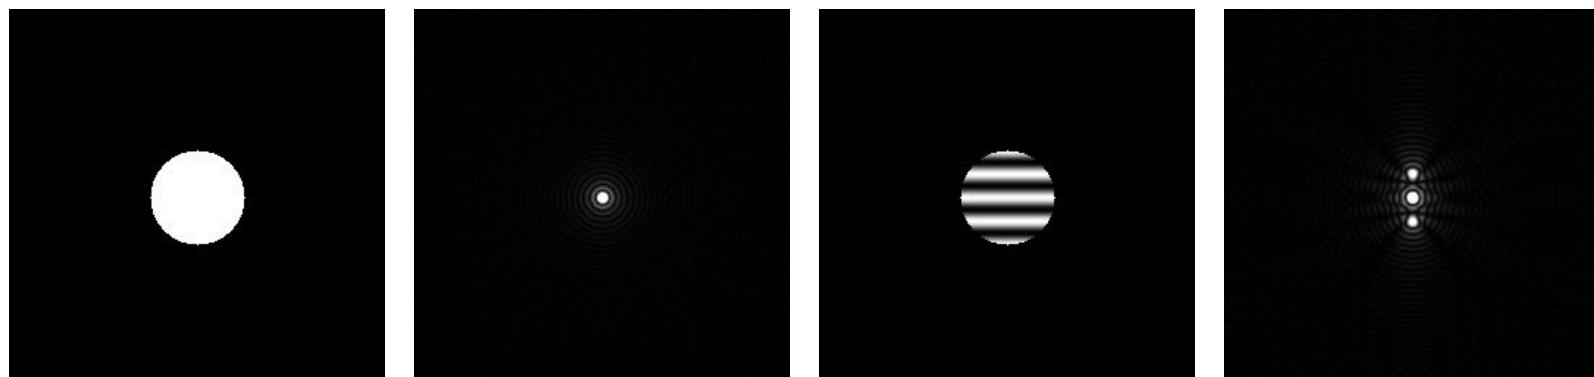


Property: Modulation

- If the original function is multiplied with an exponential like the one below and transformed, it will result in a shift of the origin of the frequency plane to (u_0, v_0) .

$$\mathcal{F} \left[f(x, y) e^{2\pi i (\frac{xu_0}{M} + \frac{yv_0}{N})} \right] = \hat{f}(u - u_0, v - v_0)$$

From left: Original image, magnitude of the Fourier spectrum, original multiplied by $1 + \cos(\omega y)$ at a relative frequency of 16, magnitude of the Fourier spectrum.



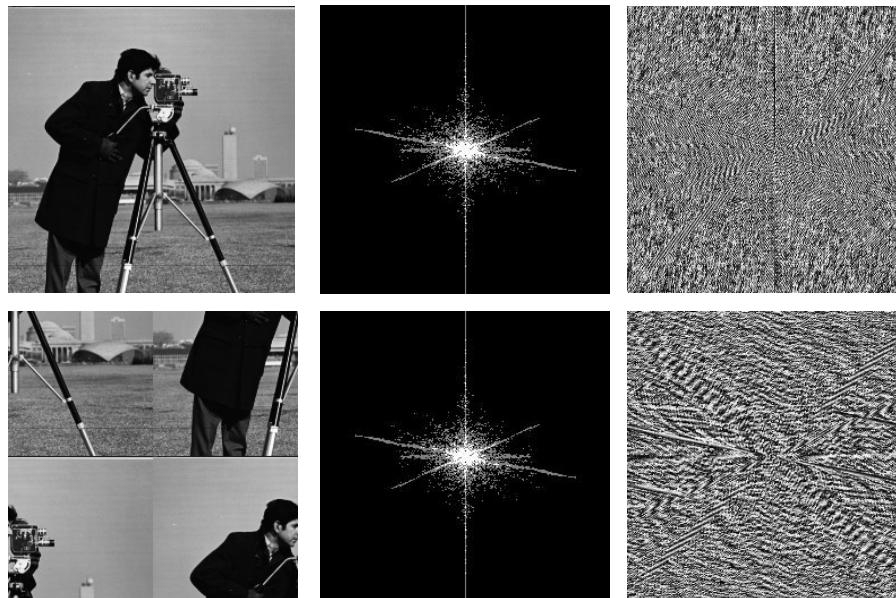
$$\text{Note : } 1 + \cos \omega y = 1 + \frac{1}{2}e^{i\omega y} + \frac{1}{2}e^{-i\omega y}.$$



Property: Translation

- (vice versa) If the image is moved, the Fourier spectrum undergoes a phase shift, but magnitude of the spectrum remains the same.

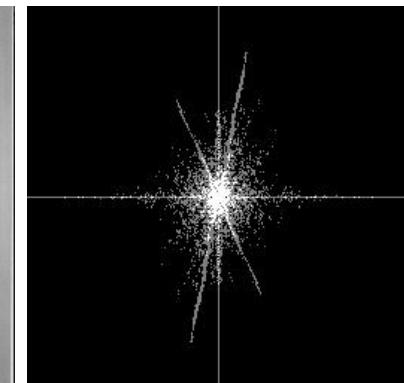
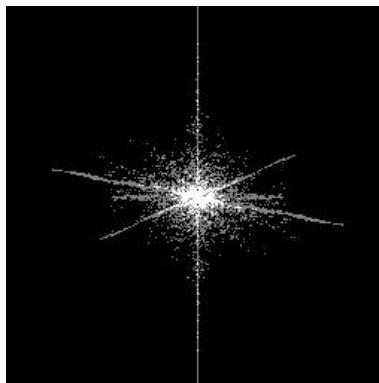
$$\mathcal{F}[f(x - x_0, y - y_0)] = \hat{f}(u, v)e^{-2\pi i(\frac{x_0 u}{M} + \frac{y_0 v}{N})}, \quad |\hat{f}(u, v)e^{-2\pi i(\frac{x_0 u}{M} + \frac{y_0 v}{N})}| = |\hat{f}(u, v)|$$





Property: Rotation

- Rotation of the original image f rotates \hat{f} by the same angle.



- Exercise: Introduce polar coordinates and perform direct substitution.
- Note: The rotation angle and phase angle are two very different kinds of angles.



Property: Scaling

- Compression (scale down) in spatial domain is same as expansion (scale up) in Fourier domain (and vice versa).

$$g(x, y) = f(S_x x, S_y y)$$

$$\hat{g}(u, v) = \frac{1}{|S_x S_y|} \hat{f}\left(\frac{u}{S_x}, \frac{v}{S_y}\right)$$





Property: Periodicity

- The DFT and its inverse are periodic with period N , for an $N \times N$ image. This means:

$$\hat{f}(u, v) = \hat{f}(u + N, v) = \hat{f}(u, v + N) = \hat{f}(u + N, v + N)$$

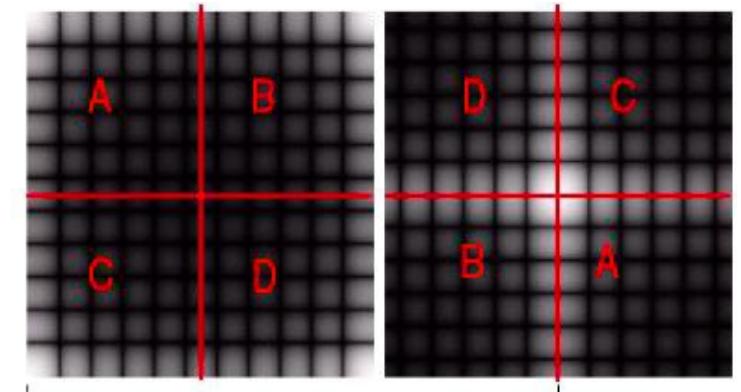
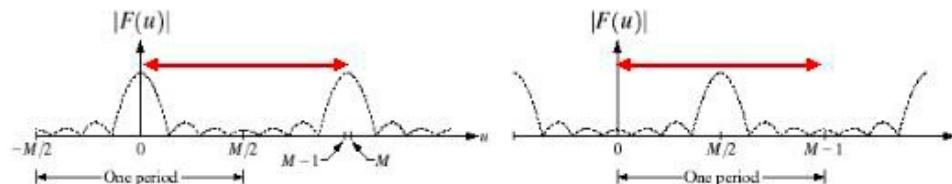
- This is easy to understand since

$$e^{-2\pi i(\frac{xu}{N})} = e^{-2\pi i(\frac{x(u+kN)}{N})}$$

- This property defines the implied symmetry in the Fourier spectrum as well as that $\hat{f}(u, v)$ repeats itself infinitely.
- However, only one period is enough to reconstruct the original image $f(x, y)$.

Property: Conjugate Symmetry

- The Fourier transform satisfies $\hat{f}(u, v) = \hat{f}^*(-u, -v)$ and $|\hat{f}(u, v)| = |\hat{f}^*(-u, -v)|$.
- With periodicity and above, we have that $\hat{f}(u, v)$ has period N and is (conjugate) symmetric around the origin.



- Thus we don't need $2N^2$ (N^2 real and N^2 imaginary) values to represent a $N \times N$ image in Fourier domain, but N^2 due to the symmetry.



Summary of good questions

- What properties does a linear shift-invariant filter have?
- How do you define a convolution?
- Why are convolutions important in linear filtering?
- How do you define a 2D Fourier transform?
- If you apply a Fourier transform to an image, what do you get?
- What information does the phase contain? And the magnitude?
- What is the Fourier transform of a convolution? Why important?
- What does separability of filters mean?
- How do you interpret a point in the Fourier domain in the spatial domain?
- How do you apply a discrete Fourier transform?
- What happens to the Fourier transform, if you translate or rotate an image?
- In what sense is the Fourier transform symmetric?



Recommended reading

- Gonzalez and Woods: Chapters 3.4, 4.3-4.7
- Szeliski: Chapter 3.2, 3.4
- Introduction to Lab 1